


Artificial Intelligence (AI)



"A number of definitions of artificial intelligence (AI) have surfaced over the last few decades, John McCarthy offers the following definition in this 2004 paper (PDF, 106 KB) (link resides outside IBM), It is the science and engineering of making intelligent machines, especially intelligent computer programs. It is related to the similar task of using computers to understand human intelligence, but AI does not have to confine itself to methods that are biologically observable. However, decades before this definition, the birth of the artificial intelligence conversation was denoted by Alan Turing seminal work, Computing Machinery and Intelligence, which was published in 1950. In this paper, Turing, often referred to as the father of computer science asks the following question, Can machines think? From there, he offers a test, now famously known as the Turing Test where a human interrogator would try to distinguish between a computer and human text response. While this test has undergone much scrutiny since its publish, it remains an important part of the history of AI as well as an ongoing concept within philosophy as it utilizes ideas around linguistics. Stuart Russell and Peter Norvig then proceeded to publish, Artificial Intelligence: A Modern Approach (link resides outside IBM), becoming one of the leading textbooks in the study of AI. In it, they delve into four potential goals or definitions of AI, which differentiates computer systems on the basis of rationality and thinking vs. acting: Human approach: Systems that think like humans Systems that act like humans Ideal approach: Systems that think rationally Systems that act rationally Alan Turing's definition would have fallen under the category of "systems that act like humans." At its simplest form, artificial intelligence is a field, which combines computer science and robust datasets, to enable problem-solving. It also encompasses sub-fields of machine learning and deep learning, which are frequently mentioned in conjunction with artificial intelligence. These disciplines are comprised of AI algorithms which seek to create expert systems which make predictions or classifications based on input data. Today, a lot of hype still surrounds AI development, which is expected of any new emerging technology in the market. As noted in Gartner's hype cycle (link resides outside IBM), product innovations like, self-driving cars and personal assistants, follow "a typical progression of innovation, from overenthusiasm through a period of disillusionment to an eventual understanding of the innovation's relevance and role in a market or domain." As Lex Fridman notes here (link resides outside IBM) in his MIT lecture in 2019, we are at the peak of inflated expectations, approaching the trough of disillusionment. As conversations emerge around the ethics of AI, we can begin to see the initial glimpses of the trough of disillusionment. To read more on where IBM stands within the conversation around AI ethics. [read more](#)"

Types of artificial intelligence-weak AI vs. strong AI




Weak AI—also called Narrow AI or Artificial Narrow Intelligence (ANI)—is AI trained and focused to perform specific tasks. Weak AI drives most of the AI that surrounds us today. ‘Narrow’ might be a more accurate descriptor for this type of AI as it is anything but weak; it enables some very robust applications, such as Apple's Siri, Amazon's Alexa, IBM Watson, and autonomous vehicles. Strong AI is made up of Artificial General Intelligence (AGI) and Artificial Super Intelligence (ASI). Artificial general intelligence (AGI), or general AI, is a theoretical form of AI where a machine would have an intelligence equaled to humans; it would have a self-aware consciousness that has the ability to solve problems, learn, and plan for the future. Artificial Super Intelligence (ASI)—also known as superintelligence—would surpass the intelligence and ability of the human brain. While strong AI is still entirely theoretical with no practical examples in use today, that doesn't mean AI researchers aren't also exploring its development. In the meantime, the best examples of ASI might be from science fiction, such as HAL, the superhuman, rogue computer assistant in 2001: A Space Odyssey.

Deep learning vs. machine learning



Since deep learning and machine learning tend to be used interchangeably, it's worth noting the nuances between the two. As mentioned above, both deep learning and machine learning are sub-fields of artificial intelligence, and deep learning is actually a sub-field of machine learning. Deep learning is actually comprised of neural networks. “Deep” in deep learning refers to a neural network comprised of more than three layers—which would be inclusive of the inputs and the output—can be considered a deep learning algorithm. This is generally represented using the following diagram.



The way in which deep learning and machine learning differ is in how each algorithm learns. Deep learning automates much of the feature extraction piece of the process, eliminating some of the manual human intervention required and enabling the use of larger data sets. You can think of deep learning as "scalable machine learning" as Lex Fridman noted in same MIT lecture from above. Classical, or "non-deep", machine learning is more dependent on human intervention to learn. Human experts determine the hierarchy of features to understand the differences between data inputs, usually requiring more structured data to learn. "Deep" machine learning can leverage labeled datasets, also known as supervised learning, to inform its algorithm, but it doesn't necessarily require a labeled dataset. It can ingest unstructured data in its raw form (e.g. text, images), and it can automatically determine the hierarchy of features which distinguish different categories of data from one another. Unlike machine learning, it doesn't require human intervention to process data, allowing us to scale machine learning in more interesting ways.

Artificial intelligence applications

There are numerous, real-world applications of AI systems today. Below are some of the most common examples: Speech Recognition: It is also known as automatic speech recognition (ASR), computer speech recognition, or speech-to-text, and it is a capability which uses natural language processing (NLP) to process human speech into a written format. Many mobile devices incorporate speech recognition into their systems to conduct voice search—e.g. Siri—or provide more accessibility around texting. Customer Service: Online chatbots are replacing human agents along the customer journey. They answer frequently asked questions (FAQs) around topics, like shipping, or provide personalized advice, cross-selling products or suggesting sizes for users, changing the way we think about customer engagement across websites and social media platforms. Examples include messaging bots on e-commerce sites with virtual agents, messaging apps, such as Slack and Facebook Messenger, and tasks usually done by virtual assistants and voice assistants. Computer Vision: This AI technology enables computers and systems to derive meaningful information from digital images, videos and other visual inputs, and based on those inputs, it can take action. This ability to provide recommendations distinguishes it from image recognition tasks. Powered by convolutional neural networks, computer vision has applications within photo tagging in social media, radiology imaging in healthcare, and self-driving cars within the automotive industry. Recommendation Engines: Using past consumption behavior data, AI algorithms can help to discover data trends that can be used to develop more effective cross-selling strategies. This is used to make relevant add-on recommendations to customers during the checkout process for online retailers. Automated stock trading: Designed to optimize stock portfolios, AI-driven high-frequency trading platforms make thousands or even millions of trades per day without human intervention.

Name: [Ahmed Osama Abdelaziz Abdelfattah](#)

Source Code

```
<html>
```

```
<head>
```

```
  <meta charset="UTF-8">
```

```
  <meta http-equiv="X-UA-Compatible" content="IE=edge">
```

```
  <meta name="viewport" content="name: احمد اسامة عبد العزيز عبد الفتاح" />
```

```
  <title> AI project</title>
```

```
</head>
```

```
<body>
```

<h1>

<font color=" black,red,yellow,pink, blue and white " size="5
>Artificial Intelligence (AI)

<p>

</p>

<p>

"A number of definitions of artificial intelligence (AI) have surfaced over the last few decades, John McCarthy offers the following definition in this 2004 paper (PDF, 106 KB) (link resides outside IBM), It is the science and engineering

of making intelligent machines, especially intelligent computer programs. It is related to the similar task of using computers to understand human intelligence, but AI does not have to confine itself to methods that are biologically observable.

However, decades before this definition, the birth of the artificial intelligence conversation was denoted by Alan Turing seminal work, Computing Machinery and Intelligence , which was published in 1950. In this paper, Turing, often referred

to as the father of computer science asks the following question, Can machines think? From there, he offers a test, now famously known as the Turing Test where a human interrogator would try to distinguish between a computer and human

text response. While this test has undergone much scrutiny since its publish, it remains an important part of the history of AI as well as an ongoing concept within philosophy as it utilizes ideas around linguistics. Stuart Russell and

Peter Norvig then proceeded to publish, Artificial Intelligence: A Modern Approach (link resides outside IBM), becoming one of the leading textbooks in the study of AI. In it, they delve into four potential goals or definitions of AI,

which differentiates computer systems on the basis of rationality and thinking vs. acting: Human approach: Systems that think like humans Systems

that act like humans Ideal approach: Systems that think rationally Systems that act rationally

Alan Turing's definition would have fallen under the category of "systems that act like humans." At its simplest form, artificial intelligence is a field, which combines computer science and robust datasets, to enable problem-solving.

It also encompasses sub-fields of machine learning and deep learning, which are frequently mentioned in conjunction with artificial intelligence. These disciplines are comprised of AI algorithms which seek to create expert systems which

make predictions or classifications based on input data. Today, a lot of hype still surrounds AI development, which is expected of any new emerging technology in the market. As noted in Gartner's hype cycle (link resides outside IBM),

product innovations like, self-driving cars and personal assistants, follow "a typical progression of innovation, from overenthusiasm through a period of disillusionment to an eventual understanding of the innovation's relevance and role

in a market or domain." As Lex Fridman notes here (link resides outside IBM) in his MIT lecture in 2019, we are at the peak of inflated expectations, approaching the trough of disillusionment. As conversations emerge around the ethics

of AI, we can begin to see the initial glimpses of the trough of disillusionment. To read more on where IBM stands within the conversation around AI ethics. [read more](https://www.ibm.com/artificial-intelligence/ethics)

` Types of artificial intelligence-weak AI
vs. strong AI `

`<p>`

``

``

Weak AI—also called Narrow AI or Artificial Narrow Intelligence (ANI)—is AI trained and focused to perform specific tasks. Weak AI drives most of the AI that surrounds us today. ‘Narrow’ might be a more accurate descriptor for this type of AI as it is

anything but weak; it enables some very robust applications, such as Apple 's Siri, Amazon 's Alexa, IBM Watson, and autonomous vehicles. Strong AI is made up of Artificial General Intelligence (AGI) and Artificial Super Intelligence (ASI).

Artificial general intelligence (AGI), or general AI, is a theoretical form of AI where a machine would have an intelligence equaled to humans; it would have a self-aware consciousness that has the ability to solve problems, learn, and

plan for the future. Artificial Super Intelligence (ASI)—also known as superintelligence—would surpass the intelligence and ability of the human brain. While strong AI is still entirely theoretical with no practical examples in use today,

that doesn 't mean AI researchers aren 't also exploring its development. In the meantime, the best examples of ASI might be from science fiction, such as HAL, the superhuman, rogue computer assistant in 2001: A Space Odyssey. ``

`</p>`

`</h2>`

`<h3>`

Deep learning vs. machine learning

<p>

Since deep learning and machine learning tend to be used interchangeably, it's worth noting the nuances between the two. As mentioned above, both deep learning and machine learning are sub-fields of artificial intelligence, and deep learning is actually

a sub-field of machine learning, Deep learning is actually comprised of neural networks. "Deep" in deep learning refers to a neural network comprised of more than three layers—which would be inclusive of the inputs and the output—can be

considered a deep learning algorithm. This is generally represented using the following diagram.

<P>

The way in which deep learning and machine learning differ is in how each algorithm learns. Deep learning automates much of the feature extraction piece of the process, eliminating some of the manual human intervention required and

enabling the use of larger data sets. You can think of deep learning as "scalable machine learning" as Lex Fridman noted in same MIT lecture from above. Classical, or "non-deep", machine learning is more dependent on human intervention

to learn. Human experts determine the hierarchy of features to understand the differences between data inputs, usually requiring more structured data to learn. "Deep" machine learning can leverage labeled datasets, also known as

supervised learning, to inform its algorithm, but it doesn't necessarily require a labeled dataset. It can ingest unstructured data in its raw form (e.g. text, images), and it can automatically determine the hierarchy of features

which distinguish different categories of data from one another. Unlike machine learning, it doesn't require human intervention to process data, allowing us to scale machine learning in more interesting ways.

</P>

</p>

</h3>

<h4>

Artificial intelligence applications

<p>

 There are numerous, real-world applications of AI systems today. Below are some of the most common examples:

<u>Speech Recognition</u>: It is also known as automatic speech recognition (ASR), computer speech recognition, or speech-to-text, and it is a capability which uses natural language processing (NLP) to process human speech into a written

format. Many mobile devices incorporate speech recognition into their systems to conduct voice search—e.g. Siri—or provide more accessibility around texting.

<u>Customer Service</u>: Online chatbots are replacing human agents along the customer journey. They answer frequently asked questions (FAQs) around topics, like shipping, or provide personalized advice, cross-selling products or suggesting

sizes for users, changing the way we think about customer engagement across websites and social media platforms. Examples include messaging bots on e-commerce sites with virtual agents, messaging apps, such as Slack and Facebook Messenger,

and tasks usually done by virtual assistants and voice assistants. Computer Vision: This AI technology enables computers and systems to derive meaningful information from digital images, videos and other visual inputs, and based

on those inputs, it can take action. This ability to provide recommendations distinguishes it from image recognition tasks. Powered by convolutional neural networks, computer vision has applications within photo tagging in social media,

radiology imaging in healthcare, and self-driving cars within the automotive industry.Recommendation Engines : Using past consumption behavior data, AI algorithms can help to discover data trends that can be used to develop more

effective cross-selling strategies. This is used to make relevant add-on recommendations to customers during the checkout process for online retailers. Automated stock trading: Designed to optimize stock portfolios, AI-driven high-frequency

trading platforms make thousands or even millions of trades per day without human intervention.

Name: Ahmed Osama Abdelaziz AbdElfatah

</html>